Kvasir-SEG Cross-Dataset Summary

Model: M2b ImageNet | Seeds: 0, 1, 2 | Resolution: 256x256 | Threshold: 0.5

Using the aggregated results from the Kvasir external-test runs, we evaluated external domain generalisation with no retraining using the same M2b ImageNet model family at 256x256 resolution and the paper-consistent threshold of 0.5. The main checkpoint choice, **valDice**, achieved strong and stable cross-seed performance with Dice 0.8875+-0.0023, IoU 0.7977+-0.0038, Precision 0.9106+-0.0101, Recall 0.8656+-0.0088, F2 0.8742+-0.0056, and area_ratio 0.9508+-0.0196. In contrast, **SeqVal-constrained** increased Recall to 0.9115+-0.0155 and F2 to 0.8924+-0.0055, but reduced Dice, IoU, and Precision and increased area_ratio to 1.1069+-0.0623, indicating a more recall-oriented operating point with moderate mask inflation. Therefore, for the main thesis narrative, the Kvasir cross-dataset experiment supports the claim that the paper's main M2b ImageNet model generalises well to an unseen external dataset without retraining, while SeqVal-constrained should be interpreted as a recall-versus-calibration trade-off rather than uniformly superior external performance.

lection Strategy	Dice	IoU	Precision	Recall	F2	area_ratio	Interpretation
Dice	0.8875+-0.0023	0.7977+-0.0038	0.9106+-0.0101	0.8656+-0.0088	0.8742+-0.0056	0.9508+-0.0196	Best balanced result; main paper-aligr
qVal-constrained	0.8655+-0.0118	0.7630+-0.0183	0.8247+-0.0323	0.9115+-0.0155	0.8924+-0.0055	1.1069+-0.0623	Higher recall/F2, but lower Dice/precis

Interpretation: valDice is the better main external result, while SeqVal-constrained is the better recall-oriented external result.